

```

import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix

# Load Fashion MNIST
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

# Reduce size for Colab (optimized even more)
x_train, y_train = x_train[:3000], y_train[:3000]
x_test, y_test = x_test[:500], y_test[:500]

# Show 10 sample images
class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
               'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
plt.figure(figsize=(10,4))
for i in range(10):
    plt.subplot(2, 5, i+1)
    plt.imshow(x_train[i], cmap='gray')
    plt.title(class_names[y_train[i]])
    plt.axis('off')
plt.tight_layout()
plt.show()

# Resize to (128x128) to save memory and convert grayscale to RGB
def preprocess_images(images):
    images = tf.expand_dims(images, -1) # Add channel dim
    images = tf.image.resize(images, [128, 128]) # Reduce from 224 to 128
    images = tf.image.grayscale_to_rgb(images)
    images = images / 255.0 # Normalize
    return images.numpy()

x_train = preprocess_images(x_train)
x_test = preprocess_images(x_test)

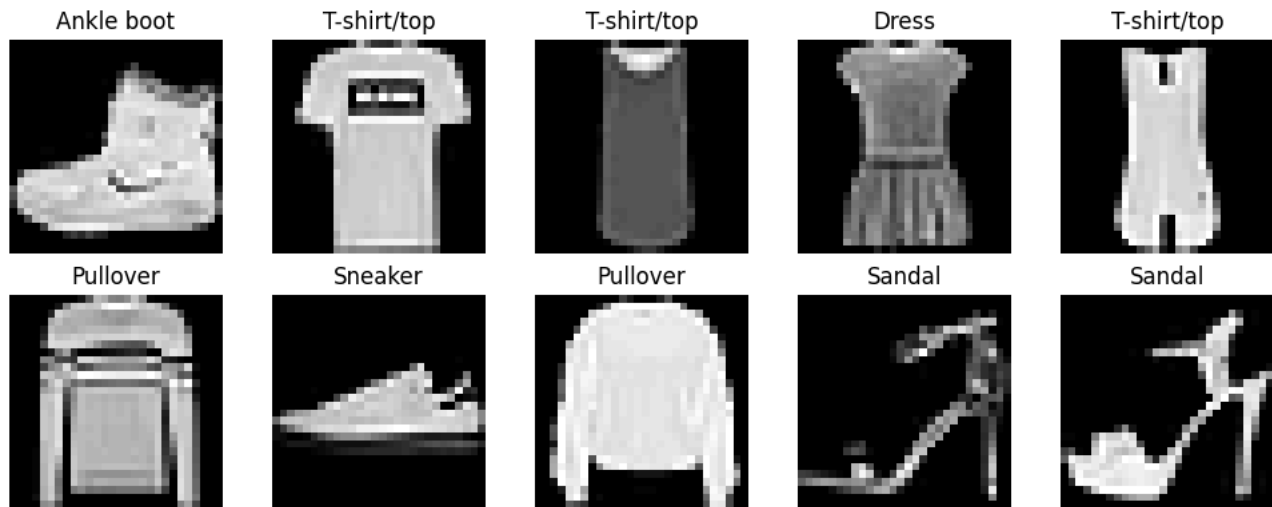
# Validation split
x_val, y_val = x_train[-300:], y_train[-300:]
x_train, y_train = x_train[:-300], y_train[:-300]

# One-hot encode labels
y_train = to_categorical(y_train, 10)
y_val = to_categorical(y_val, 10)
y_test = to_categorical(y_test, 10)

# Data Augmentation (lighter)
datagen = ImageDataGenerator(rotation_range=10, zoom_range=0.05, horizontal_flip=True)
datagen.fit(x_train)

```

 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>
 29515/29515 — 0s 0us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>
 26421880/26421880 — 2s 0us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>
 5148/5148 — 0s 1us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>
 4422102/4422102 — 1s 0us/step



```

import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns

# Load dataset
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

# Reduce size for Colab
x_train, y_train = x_train[:5000], y_train[:5000]
x_test, y_test = x_test[:1000], y_test[:1000]

# Resize and convert to RGB
x_train = tf.image.resize_with_pad(tf.expand_dims(x_train, -1), 128, 128)
x_train = tf.image.grayscale_to_rgb(x_train)
x_test = tf.image.resize_with_pad(tf.expand_dims(x_test, -1), 128, 128)
x_test = tf.image.grayscale_to_rgb(x_test)

# Normalize
x_train, x_test = x_train / 255.0, x_test / 255.0
x_train, x_test = x_train.numpy(), x_test.numpy()

# Validation split
x_val, y_val = x_train[-500:], y_train[-500:]
x_train, y_train = x_train[:-500], y_train[:-500]

# One-hot encode
y_train = to_categorical(y_train, 10)
y_val = to_categorical(y_val, 10)
y_test = to_categorical(y_test, 10)

# Data Augmentation
datagen = ImageDataGenerator(rotation_range=15, zoom_range=0.1, horizontal_flip=True)
datagen.fit(x_train)

# Load MobileNetV2
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam

base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(128, 128, 3))

# Freeze base layers
for layer in base_model.layers:
    layer.trainable = False

# Add custom head

```

```
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.5)(x)
x = Dense(256, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

model = Model(inputs=base_model.input, outputs=predictions)

# Compile model
model.compile(optimizer=Adam(learning_rate=0.0001),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train model
history = model.fit(datagen.flow(x_train, y_train, batch_size=32),
                    validation_data=(x_val, y_val),
                    epochs=10)

# -----
# Feature Map Visualization
# -----
from tensorflow.keras.models import Model

# Choose a conv layer
layer_name = 'block_1_expand_relu' # You can change this layer name
feature_model = Model(inputs=model.input, outputs=model.get_layer(layer_name).output)

# Get feature maps for one image
sample_image = x_train[0:1]
feature_maps = feature_model.predict(sample_image)

# Plot first 8 feature maps
plt.figure(figsize=(12, 6))
for i in range(8):
    plt.subplot(2, 4, i + 1)
    plt.imshow(feature_maps[0, :, :, i], cmap='viridis')
    plt.axis('off')
plt.suptitle(f'Feature Maps from layer: {layer_name}')
plt.tight_layout()
plt.show()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_9406464/9406464 2s 0us/step

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class has a `warn_if_super_not_called` attribute, but it is not called. This may lead to unexpected behavior.
self._warn_if_super_not_called()

Epoch 1/10
141/141 ————— 39s 199ms/step - accuracy: 0.3187 - loss: 2.0435 - val_accuracy: 0.7960 - val_loss: 0.6196

Epoch 2/10
141/141 ————— 18s 128ms/step - accuracy: 0.7137 - loss: 0.8882 - val_accuracy: 0.8520 - val_loss: 0.4400

Epoch 3/10
141/141 ————— 17s 121ms/step - accuracy: 0.7464 - loss: 0.7284 - val_accuracy: 0.8540 - val_loss: 0.3849

Epoch 4/10
141/141 ————— 18s 128ms/step - accuracy: 0.7545 - loss: 0.6619 - val_accuracy: 0.8480 - val_loss: 0.3832

Epoch 5/10
141/141 ————— 17s 123ms/step - accuracy: 0.7707 - loss: 0.6292 - val_accuracy: 0.8780 - val_loss: 0.3580

Epoch 6/10
141/141 ————— 18s 128ms/step - accuracy: 0.8009 - loss: 0.5509 - val_accuracy: 0.8800 - val_loss: 0.3339

Epoch 7/10
141/141 ————— 19s 120ms/step - accuracy: 0.8079 - loss: 0.5455 - val_accuracy: 0.8820 - val_loss: 0.3219


Epoch 8/10
141/141 ————— 21s 127ms/step - accuracy: 0.8209 - loss: 0.5019 - val_accuracy: 0.8940 - val_loss: 0.3090

Epoch 9/10
141/141 ————— 17s 120ms/step - accuracy: 0.8131 - loss: 0.4945 - val_accuracy: 0.8880 - val_loss: 0.3085

Epoch 10/10
141/141 ————— 18s 125ms/step - accuracy: 0.8212 - loss: 0.5035 - val_accuracy: 0.8920 - val_loss: 0.3041

1/1 ————— 1s 525ms/step

Feature Maps from layer: block_1_expand_relu



```
# =====EVALUATE ON TEST DATA =====
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=1)
print(f"\nTest Accuracy: {test_acc:.4f}")
print(f"Test Loss: {test_loss:.4f}")
```

```
# ==PREDICTIONS ==
y_pred_probs = model.predict(x_test)
y_pred = np.argmax(y_pred_probs, axis=1)
y_true = np.argmax(y_test, axis=1)
```

```
# == CLASSIFICATION REPORT ==
print("\nClassification Report:\n")
print(classification_report(y_true, y_pred, target_names=[
    "T-shirt/top", "Trouser", "Pullover", "Dress", "Coat",
    "Sandal", "Shirt", "Sneaker", "Bag", "Ankle boot"
]))
```

```
# == CONFUSION MATRIX ==
conf_matrix = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
            xticklabels=[
                "T-shirt", "Trouser", "Pullover", "Dress", "Coat",
                "Sandal", "Shirt", "Sneaker", "Bag", "Boot"
            ],
            yticklabels=[
                "T-shirt", "Trouser", "Pullover", "Dress", "Coat",
                "Sandal", "Shirt", "Sneaker", "Bag", "Boot"
            ],
            cbar_kws={'label': 'Confusion Matrix'})
```

```

y_test_names = ['T-shirt', 'Trouser', 'Pullover', 'Dress', 'Coat',
                'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Boot']

plt.title("Confusion Matrix")
plt.ylabel("Actual")
plt.xlabel("Predicted")
plt.show()

# == ACCURACY & LOSS PLOTS ==
plt.figure(figsize=(14, 5))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()

```

32/32 ————— 8s 270ms/step - accuracy: 0.8560 - loss: 0.3812

Test Accuracy: 0.8620

Test Loss: 0.3882

32/32 ————— 9s 111ms/step

Classification Report:

	precision	recall	f1-score	support
T-shirt/top	0.81	0.85	0.83	107
Trouser	1.00	0.98	0.99	105
Pullover	0.76	0.86	0.81	111
Dress	0.88	0.84	0.86	93
Coat	0.82	0.81	0.81	115
Sandal	0.97	0.86	0.91	87
Shirt	0.66	0.61	0.63	97
Sneaker	0.86	0.91	0.88	95
Bag	0.99	0.96	0.97	95
Ankle boot	0.92	0.96	0.94	95
accuracy			0.86	1000
macro avg	0.87	0.86	0.86	1000
weighted avg	0.86	0.86	0.86	1000

